

Assignment is at the bottom!

```
In [2]: from sklearn.linear_model import LogisticRegression
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np


from pylab import rcParams
rcParams['figure.figsize'] = 20, 10

from sklearn.linear_model import LogisticRegression as Model
```

```
In [3]: y = np.concatenate([np.zeros(10), np.ones(10)])
x = np.linspace(0, 10, len(y))
```

```
In [4]: plt.scatter(x, y, c=y)
```

```
Out[4]: <matplotlib.collections.PathCollection at 0x7f8d695c1870>
```




```
In [5]: model = LogisticRegression()
```

```
In [6]: model.fit(x.reshape(-1, 1),y)
```

```
Out[6]: ▼ LogisticRegression  
LogisticRegression()
```

```
In [7]: plt.scatter(x,y, c=y)  
plt.plot(x, model.predict_proba(x.reshape(-1, 1))[:,1])
```

```
Out[7]: [matplotlib.lines.Line2D at 0x7f8d69ea28c0>]
```




```
In [8]: b, b0 = model.coef_, model.intercept_  
model.coef_, model.intercept_
```

```
Out[8]: (array([[1.46709085]]), array([-7.33542562]))
```

```
In [9]: plt.plot(x, 1/(1+np.exp(-x)))
```

```
Out[9]: [<matplotlib.lines.Line2D at 0x7f8d6a947ac0>]
```




```
In [10]: b
```

```
Out[10]: array([[1.46709085]])
```

```
In [11]: plt.plot(x, 1/(1+np.exp(-(b[0]*x + b0))))
```

```
Out[11]: [<matplotlib.lines.Line2D at 0x7f8d6a9d10f0>]
```



```
In [12]: from mpl_toolkits.mplot3d import Axes3D # noqa: F401 unused import

import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator, FormatStrFormatter
import numpy as np

fig = plt.figure()
ax = fig.gca(projection='3d')# Doesn't like "projection"

# Make data.
X = np.arange(-10, 10, 0.25)
Y = np.arange(-10, 10, 0.25)
X, Y = np.meshgrid(X, Y)
R = np.sqrt(X**2 + Y**2)
Z = 1/(1+np.exp(-(b[0]*X + b[0]*Y + b0)))
surf = ax.plot_surface(X, Y, Z, cmap=cm.coolwarm,
                      linewidth=0, antialiased=False)
```

```
-----
TypeError                                Traceback (most recent call
last)
Cell In[12], line 10
      6 import numpy as np
      9 fig = plt.figure()
--> 10 ax = fig.gca(projection='3d')
      12 # Make data.
      13 X = np.arange(-10, 10, 0.25)

TypeError: FigureBase.gca() got an unexpected keyword argument 'proje
ction'
```

<Figure size 2000x1000 with 0 Axes>

```
In [13]: X
```

```
-----
NameError                                Traceback (most recent call
last)
Cell In[13], line 1
--> 1 X

NameError: name 'X' is not defined
```

In [13]: Y


```
Out[13]: array([[ -10.   ,  -10.   ,  -10.   , ...,  -10.   ,  -10.   ,  -10.   ],
                [  -9.75,  -9.75,  -9.75, ...,  -9.75,  -9.75,  -9.75],
                [  -9.5 ,  -9.5 ,  -9.5 , ...,  -9.5 ,  -9.5 ,  -9.5 ],
                ...,
                [   9.25,   9.25,   9.25, ...,   9.25,   9.25,   9.25],
                [   9.5 ,   9.5 ,   9.5 , ...,   9.5 ,   9.5 ,   9.5 ],
                [   9.75,   9.75,   9.75, ...,   9.75,   9.75,   9.75]])
```

What if the data doesn't really fit this pattern?

```
In [14]: y = np.concatenate([np.zeros(10), np.ones(10), np.zeros(10)])
        x = np.linspace(0, 10, len(y))
```

```
In [15]: plt.scatter(x,y, c=y)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x7f8d6b518670>
```




```
In [16]: model.fit(x.reshape(-1, 1), y)
```

```
Out[16]: ▼ LogisticRegression
          LogisticRegression()
```

```
In [17]: plt.scatter(x,y)
plt.plot(x, model.predict_proba(x.reshape(-1, 1)))
```

```
Out[17]: [<matplotlib.lines.Line2D at 0x7f8d6b533e20>,
<matplotlib.lines.Line2D at 0x7f8d6b533e80>]
```



```
In [18]: model1 = LogisticRegression()
model1.fit(x[:15].reshape(-1, 1),y[:15])
```


```
Out[18]: ▼ LogisticRegression
LogisticRegression()
```

```
In [19]: model2 = LogisticRegression()
model2.fit(x[15:].reshape(-1, 1),y[15:])
```

```
Out[19]: ▼ LogisticRegression
LogisticRegression()
```

```
In [20]: plt.scatter(x,y, c=y)
plt.plot(x, model1.predict_proba(x.reshape(-1, 1))[:,1] * model2.predi
```

```
Out[20]: [<matplotlib.lines.Line2D at 0x7f8d6c2ab910>]
```



```
In [21]: df = pd.read_csv('../data/adult.data', index_col=False)
golden = pd.read_csv('../data/adult.test', index_col=False)
```

```
In [22]: from sklearn import preprocessing
enc = preprocessing.OrdinalEncoder()
```

```
In [23]: transform_columns = ['sex', 'workclass', 'education', 'marital-status',
                             'occupation', 'relationship', 'race', 'sex',
                             'native-country', 'salary']
```

```
In [24]: x = df.copy()

x[transform_columns] = enc.fit_transform(df[transform_columns])

golden['salary'] = golden.salary.replace(' <=50K.', ' <=50K').replace(
xt = golden.copy()

xt[transform_columns] = enc.transform(golden[transform_columns])
```

```
In [25]: df.salary.unique()
```

```
Out[25]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [26]: golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K').
```

```
Out[26]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [27]: model.fit(preprocessing.scale(x.drop('salary', axis=1)), x.salary)
```

```
Out[27]: ▼ LogisticRegression
LogisticRegression()
```

```
In [28]: pred = model.predict(preprocessing.scale(x.drop('salary', axis=1)))
pred_test = model.predict(preprocessing.scale(xt.drop('salary', axis=1)))
```

```
In [29]: x.head()
```

```
Out[29]:
```

	age	workclass	fnlwgt	education	education- num	marital- status	occupation	relationship	race	sex
0	39	7.0	77516	9.0	13	4.0	1.0	1.0	4.0	1.0
1	50	6.0	83311	9.0	13	2.0	4.0	0.0	4.0	1.0
2	38	4.0	215646	11.0	9	0.0	6.0	1.0	4.0	1.0
3	53	4.0	234721	1.0	7	2.0	6.0	0.0	2.0	1.0
4	28	4.0	338409	9.0	13	2.0	10.0	5.0	2.0	0.0

```
In [30]: from sklearn.metrics import (
    accuracy_score,
    classification_report,
    confusion_matrix, auc, roc_curve
)
```

```
In [31]: accuracy_score(x.salary, pred)
```

```
Out[31]: 0.8250360861152913
```

```
In [32]: confusion_matrix(x.salary, pred)
```

```
Out[32]: array([[23300, 1420],
               [ 4277, 3564]])
```



```
In [33]: print(classification_report(x.salary, pred))
```

	precision	recall	f1-score	support
0.0	0.84	0.94	0.89	24720
1.0	0.72	0.45	0.56	7841
accuracy			0.83	32561
macro avg	0.78	0.70	0.72	32561
weighted avg	0.81	0.83	0.81	32561

```
In [34]: print(classification_report(xt.salary, pred_test))
```

	precision	recall	f1-score	support
0.0	0.85	0.94	0.89	12435
1.0	0.70	0.45	0.55	3846
accuracy			0.82	16281
macro avg	0.77	0.69	0.72	16281
weighted avg	0.81	0.82	0.81	16281

Assignment

1. Use your own dataset (Heart.csv is acceptable), create a train and a test set, and build 2 models: Logistic Regression and Decision Tree (shallow). Compare the test results using classification_report and confusion_matrix. Explain which algorithm is optimal

In [48]: *# Data load*

```
heart = pd.read_csv('../data/Heart.csv', index_col=False)
heart.head()
```

Out[48]:

	Unnamed: 0	Age	Sex	ChestPain	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak
0	1	63	1	typical	145	233	1	2	150	0	2.6
1	2	67	1	asymptomatic	160	286	0	2	108	1	1.5
2	3	67	1	asymptomatic	120	229	0	2	129	1	2.6
3	4	37	1	nonanginal	130	250	0	0	187	0	3.5
4	5	41	0	nontypical	130	204	0	2	172	0	1.4

In [49]: *# Data Transform and Cleaning*

```
heart_transform_clmns = ['ChestPain', 'Thal']
heart_clean = heart.dropna()
ht = heart_clean.copy()
ht[heart_transform_clmns] = enc.fit_transform(heart_clean[heart_transform_clmns])
```

In [138]: *# Model Definition*

```
from sklearn.tree import DecisionTreeClassifier
logi = LogisticRegression()
short_tree = DecisionTreeClassifier(criterion='entropy', max_depth=2)
```

In [127]: *# Train Test Split*

```
from sklearn.model_selection import train_test_split
ht_x_train, ht_x_test, ht_y_train, ht_y_test = train_test_split(ht.drop('AHD', axis=1),
                                                                ht.AHD, test_size=0.3)
```

In [139]: *# Model Fit*

```
logi.fit(preprocessing.scale(ht_x_train.values), ht_y_train)
short_tree.fit(preprocessing.scale(ht_x_train.values), ht_y_train)
```

Out[139]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=2)
```

In [140]: *# Testing*

```
logi_predictions = logi.predict(preprocessing.scale(ht_x_test.values))
short_tree_predictions = short_tree.predict(preprocessing.scale(ht_x_test.values))
```

In [141]: *# Accuracy*

```
accuracy_score(ht_y_test, logi_predictions), accuracy_score(ht_y_test,
```

Out[141]: (0.9, 0.8833333333333333)

In [142]: *# Confusion Matrix*

```
confusion_matrix(ht_y_test, logi_predictions), confusion_matrix(ht_y_test,
```

Out[142]: (array([[27, 2],
[4, 27]]),
array([[25, 4],
[3, 28]]))

In [143]: *# Classification Report*

```
print(classification_report(ht_y_test, logi_predictions),classification_report(ht_y_test, shallow_predictions))
```

	precision	recall	f1-score	support
No	0.87	0.93	0.90	29
Yes	0.93	0.87	0.90	31
accuracy			0.90	60
macro avg	0.90	0.90	0.90	60
weighted avg	0.90	0.90	0.90	60

	precision	recall	f1-score	support
No	0.89	0.86	0.88	29
Yes	0.88	0.90	0.89	31
accuracy			0.88	60
macro avg	0.88	0.88	0.88	60
weighted avg	0.88	0.88	0.88	60

Results Disucssion:

First off I would say that the overall results regarding the accuracy, confusion matrix, and classification reports vary wildly depending on what the test train split uses for the training data. I was seeing the accuracy and subsequent report values varying from 0.50 to above 0.90. I used the current test train split as I think it is fairly representative of the general trends between the two models. Broadly speaking with this data set the logistic regression model performed roughly equal to or, slightly better than, the shallow decision tree. The trend that appeared to be holding the shallow decision tree back was overall inferior precision while generally having higher recall. Between the two models it is a close call but ultimately I would choose the logistic regression for it's more consistant performance.

2. Repeat 1. but let the Decision Tree be much deeper to allow over-fitting. Compare the two models' test results again, and explain which is optimal

In [144]: *# Model Definition*

```
tall_tree = DecisionTreeClassifier(criterion='entropy', max_depth=15)
```

In [145]: *# Model Fit*

```
tall_tree.fit(preprocessing.scale(ht_x_train.values), ht_y_train)
```

Out[145]:

```
DecisionTreeClassifier  
DecisionTreeClassifier(criterion='entropy', max_depth=15)
```

In [146]: *# Testing*

```
tall_tree_predictions = tall_tree.predict(preprocessing.scale(ht_x_test.values))
```

In [147]: *# Accuracy*

```
accuracy_score(ht_y_test, logi_predictions), accuracy_score(ht_y_test, tall_tree_predictions)
```

Out[147]: (0.9, 0.7)

In [148]: *# Confusion Matrix*

```
confusion_matrix(ht_y_test, logi_predictions), confusion_matrix(ht_y_test, tall_tree_predictions)
```

Out[148]: (array([[27, 2],
[4, 27]]),
array([[19, 10],
[8, 23]]))

In [149]: *# Classification Report*

```
print(classification_report(ht_y_test, logi_predictions),classification_report
```

	precision	recall	f1-score	support
No	0.87	0.93	0.90	29
Yes	0.93	0.87	0.90	31
accuracy			0.90	60
macro avg	0.90	0.90	0.90	60
weighted avg	0.90	0.90	0.90	60

	precision	recall	f1-score	support
No	0.70	0.66	0.68	29
Yes	0.70	0.74	0.72	31
accuracy			0.70	60
macro avg	0.70	0.70	0.70	60
weighted avg	0.70	0.70	0.70	60

Results Disucssion:

As mentioned above the actual values for the accuracy, confusion matrix, and the classification report vary significantly depending on the train test split data. The performance between the two models remained fairly consistent between the training sets. Here we can clearly see the deep decision tree overfitting for the training set. This is evident in the low accuracy values, amount of missed estimates in the confusion matrix, and the precision, recall, and f1 scores, when compared to the logistic regression model from earlier. Between the two models the logistic regression is clearly more consistently capable at prediction with this data set.