

Assignment is below at the bottom

Video 13.1 <https://www.youtube.com/watch?v=kIGHE7Cfe1s>
(<https://www.youtube.com/watch?v=kIGHE7Cfe1s>)

Video 13.2 <https://www.youtube.com/watch?v=Rm9bJcDd1KU>
(<https://www.youtube.com/watch?v=Rm9bJcDd1KU>)

Video 13.3 <https://youtu.be/6HjZk-3LsjE> (<https://youtu.be/6HjZk-3LsjE>)

```
In [11]: from keras.callbacks import TensorBoard
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: # Demo data set up

(xtrain, ytrain), (xtest, ytest) = mnist.load_data()

xtrain = xtrain.astype('float32') / 255.
xtest = xtest.astype('float32') / 255.
xtrain = xtrain.reshape((len(xtrain), np.prod(xtrain.shape[1:])))
xtest = xtest.reshape((len(xtest), np.prod(xtest.shape[1:])))
xtrain.shape, xtest.shape
```

```
Out[2]: ((60000, 784), (10000, 784))
```

In [28]:

```

# this is the size of our encoded representations
encoding_dim = 4 # 32 floats -> compression of factor 24.5, assuming

# this is our input placeholder
x = input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
x = Dense(256, activation='relu')(x)
x = Dense(128, activation='relu')(x)
encoded = Dense(encoding_dim, activation='relu')(x)

# "decoded" is the lossy reconstruction of the input
x = Dense(128, activation='relu')(encoded)
x = Dense(256, activation='relu')(x)
decoded = Dense(784, activation='sigmoid')(x)

# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)

encoder = Model(input_img, encoded)

# create a placeholder for an encoded (32-dimensional) input
encoded_input = Input(shape=(encoding_dim,))
# retrieve the last layer of the autoencoder model
dcd1 = autoencoder.layers[-1]
dcd2 = autoencoder.layers[-2]
dcd3 = autoencoder.layers[-3]

# create the decoder model
decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))

```

In [29]: `autoencoder.compile(optimizer='adam', loss='binary_crossentropy')`

```

In [ ]: autoencoder.fit(xtrain, xtrain,
                        epochs=100,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(xtest, xtest),
                        #callbacks=[TensorBoard(log_dir='/tmp/autoencoder')])

```


```
In [38]: encoded_imgs
```

```
Out[38]: array([[11.943697 ,  9.005527 , 12.027234 , 32.89881  ],
                [23.76052  , 13.926956 ,  5.6552634,  8.942506 ],
                [35.62965  , 34.729908 , 24.666973 , 41.5047   ],
                ...,
                [ 5.3135986, 11.108302 , 14.398285 , 17.106884 ],
                [ 4.376413 , 19.419018 , 15.854642 , 11.992302 ],
                [ 7.41167  , 18.699078 , 30.420742 , 11.0364065]], dtype=float
32)
```

```
In [52]: noise = np.random.normal(20,4, (4,4))
noise_preds = decoder.predict(noise)
```

```
In [55]: plt.imshow(noise_preds[1].reshape(28,28))
```

```
Out[55]: <matplotlib.image.AxesImage at 0x13bf35780>
```



```
In [41]: np.max(encoded_imgs)
```

```
Out[41]: 54.59457
```

```
In [32]: encoded_imgs = encoder.predict(xtest)
         decoded_imgs = decoder.predict(encoded_imgs)
         import matplotlib.pyplot as plt

         n = 20 # how many digits we will display
         plt.figure(figsize=(40, 4))
         for i in range(n):
             # display original
             ax = plt.subplot(2, n, i + 1)
             plt.imshow(xtest[i].reshape(28, 28))
             plt.gray()
             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)

             # display reconstruction
             ax = plt.subplot(2, n, i + 1 + n)
             plt.imshow(decoded_imgs[i].reshape(28, 28))
             plt.gray()
             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)
         plt.show()
```




```
In [33]: encoded_imgs
```

```
Out[33]: array([[11.943697 ,  9.005527 , 12.027234 , 32.89881  ],
                [23.76052  , 13.926956 ,  5.6552634,  8.942506 ],
                [35.62965  , 34.729908 , 24.666973 , 41.5047   ],
                ...,
                [ 5.3135986, 11.108302 , 14.398285 , 17.106884 ],
                [ 4.376413 , 19.419018 , 15.854642 , 11.992302 ],
                [ 7.41167  , 18.699078 , 30.420742 , 11.0364065]], dtype=float
32)
```

```
In [26]: %matplotlib inline
```


```
In [34]: plt.scatter(encoded_imgs[:,1], encoded_imgs[:,0], s=1, c=ytest, cmap='  
# plt.show())
```

Out[34]: <matplotlib.collections.PathCollection at 0x13c081978>




```
In [35]: plt.scatter(encoded_imgs[:,1], encoded_imgs[:,3], s=1, c=ytest, cmap='  
# plt.show())
```

Out[35]: <matplotlib.collections.PathCollection at 0x13b695e10>




```
In [36]: plt.scatter(encoded_imgs[:,1], encoded_imgs[:,2], s=1, c=ytest, cmap='
# plt.show())
```

Out[36]: <matplotlib.collections.PathCollection at 0x13b6eaf60>



```
In [37]: from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(encoded_imgs[:,1], encoded_imgs[:,2], encoded_imgs[:,3], c=
```

Out[37]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x13c0e7da0>



Assignment

1. change the `encoding_dim` through various values (range(2,18,2)) and save the loss you can get. Plot the 8 pairs of dimensions vs loss on a scatter plot

In [7]: *# Assignment Data Setup*

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
x_train.shape, x_test.shape
```

Out[7]: ((60000, 784), (10000, 784))

In [12]: *# Loop Setup*

```
dimensions = range(2,18,2)
losses = []

for encoding_dim in dimensions:

    # Layer Setup this seems inefficient to remake the layer for each
    x = input_img = Input(shape=(784,))
    # "encoded" is the encoded representation of the input
    x = Dense(256, activation='relu')(x)
    x = Dense(128, activation='relu')(x)
    encoded = Dense(encoding_dim, activation='relu')(x)

    # "decoded" is the lossy reconstruction of the input
    x = Dense(128, activation='relu')(encoded)
    x = Dense(256, activation='relu')(x)
    decoded = Dense(784, activation='sigmoid')(x)

    # this model maps an input to its reconstruction
    autoencoder = Model(input_img, decoded)

    encoder = Model(input_img, encoded)

    # create a placeholder for an encoded (32-dimensional) input
    encoded_input = Input(shape=(encoding_dim,))
    # retrieve the last layer of the autoencoder model
    dcd1 = autoencoder.layers[-1]
    dcd2 = autoencoder.layers[-2]
    dcd3 = autoencoder.layers[-3]

    # create the decoder model
    decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))

    autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

```

history = autoencoder.fit(x_train, x_train,
                          epochs=100,
                          batch_size=256,
                          shuffle=True,
                          verbose=0, # processes faster when not set to verbose
                          validation_data=(x_test, x_test))

plt.plot(history.history['loss'], label=f'{encoding_dim} dimension')

autoencoder.summary()

loss = autoencoder.evaluate(x_train, x_train, verbose=0)

print(loss)

losses.append(loss)

plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs for Different Encoded Dimensions')
plt.legend()
plt.show()
printout = [item for sublist in zip(dimensions, losses) for item in sublist]
print("Dimensions, Loss, Accuracy:", printout)

```

Model: "model_6"

Layer (type)	Output Shape	Param #
=====		
input_5 (InputLayer)	[(None, 784)]	0
dense_12 (Dense)	(None, 256)	200960
dense_13 (Dense)	(None, 128)	32896
dense_14 (Dense)	(None, 2)	258
dense_15 (Dense)	(None, 128)	384
dense_16 (Dense)	(None, 256)	33024
dense_17 (Dense)	(None, 784)	201488
=====		
Total params: 469010 (1.79 MB)		
Trainable params: 469010 (1.79 MB)		
Non-trainable params: 0 (0.00 Byte)		

0.16411425173282623

Model: "model_9"

Layer (type)	Output Shape	Param #
input_7 (InputLayer)	[(None, 784)]	0
dense_18 (Dense)	(None, 256)	200960
dense_19 (Dense)	(None, 128)	32896
dense_20 (Dense)	(None, 4)	516
dense_21 (Dense)	(None, 128)	640
dense_22 (Dense)	(None, 256)	33024
dense_23 (Dense)	(None, 784)	201488

=====
 Total params: 469524 (1.79 MB)
 Trainable params: 469524 (1.79 MB)
 Non-trainable params: 0 (0.00 Byte)

0.13520190119743347

Model: "model_12"

Layer (type)	Output Shape	Param #
input_9 (InputLayer)	[(None, 784)]	0
dense_24 (Dense)	(None, 256)	200960
dense_25 (Dense)	(None, 128)	32896
dense_26 (Dense)	(None, 6)	774
dense_27 (Dense)	(None, 128)	896
dense_28 (Dense)	(None, 256)	33024
dense_29 (Dense)	(None, 784)	201488

=====
 Total params: 470038 (1.79 MB)
 Trainable params: 470038 (1.79 MB)
 Non-trainable params: 0 (0.00 Byte)

0.11663084477186203

Model: "model_15"

Layer (type)	Output Shape	Param #
input_11 (InputLayer)	[(None, 784)]	0
dense_30 (Dense)	(None, 256)	200960
dense_31 (Dense)	(None, 128)	32896
dense_32 (Dense)	(None, 8)	1032
dense_33 (Dense)	(None, 128)	1152
dense_34 (Dense)	(None, 256)	33024
dense_35 (Dense)	(None, 784)	201488

```

Total params: 470552 (1.80 MB)
Trainable params: 470552 (1.80 MB)
Non-trainable params: 0 (0.00 Byte)

```

```

0.11094623804092407
Model: "model_18"

```

Layer (type)	Output Shape	Param #
input_13 (InputLayer)	[(None, 784)]	0
dense_36 (Dense)	(None, 256)	200960
dense_37 (Dense)	(None, 128)	32896
dense_38 (Dense)	(None, 10)	1290
dense_39 (Dense)	(None, 128)	1408
dense_40 (Dense)	(None, 256)	33024
dense_41 (Dense)	(None, 784)	201488

```

Total params: 471066 (1.80 MB)
Trainable params: 471066 (1.80 MB)
Non-trainable params: 0 (0.00 Byte)

```

```

0.1008804589509964
Model: "model_21"

```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

input_15 (InputLayer)	[(None, 784)]	0
dense_42 (Dense)	(None, 256)	200960
dense_43 (Dense)	(None, 128)	32896
dense_44 (Dense)	(None, 12)	1548
dense_45 (Dense)	(None, 128)	1664
dense_46 (Dense)	(None, 256)	33024
dense_47 (Dense)	(None, 784)	201488

=====
Total params: 471580 (1.80 MB)

Trainable params: 471580 (1.80 MB)

Non-trainable params: 0 (0.00 Byte)

0.09382828325033188

Model: "model_24"

Layer (type)	Output Shape	Param #
input_17 (InputLayer)	[(None, 784)]	0
dense_48 (Dense)	(None, 256)	200960
dense_49 (Dense)	(None, 128)	32896
dense_50 (Dense)	(None, 14)	1806
dense_51 (Dense)	(None, 128)	1920
dense_52 (Dense)	(None, 256)	33024
dense_53 (Dense)	(None, 784)	201488

=====
Total params: 472094 (1.80 MB)

Trainable params: 472094 (1.80 MB)

Non-trainable params: 0 (0.00 Byte)

0.0918399766087532

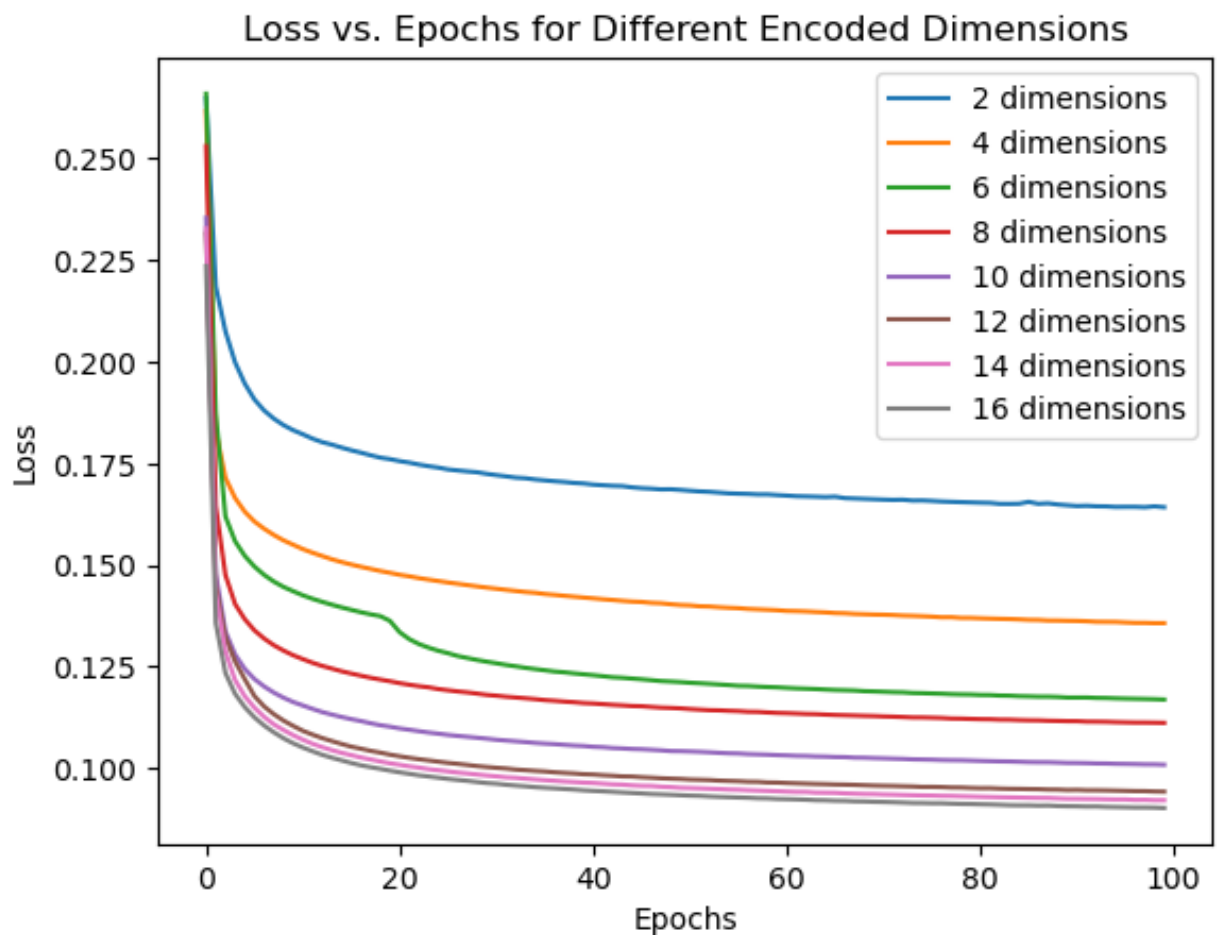
Model: "model_27"

Layer (type)	Output Shape	Param #
input_19 (InputLayer)	[(None, 784)]	0

dense_54 (Dense)	(None, 256)	200960
dense_55 (Dense)	(None, 128)	32896
dense_56 (Dense)	(None, 16)	2064
dense_57 (Dense)	(None, 128)	2176
dense_58 (Dense)	(None, 256)	33024
dense_59 (Dense)	(None, 784)	201488

```
=====
Total params: 472608 (1.80 MB)
Trainable params: 472608 (1.80 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
0.09008263051509857
```



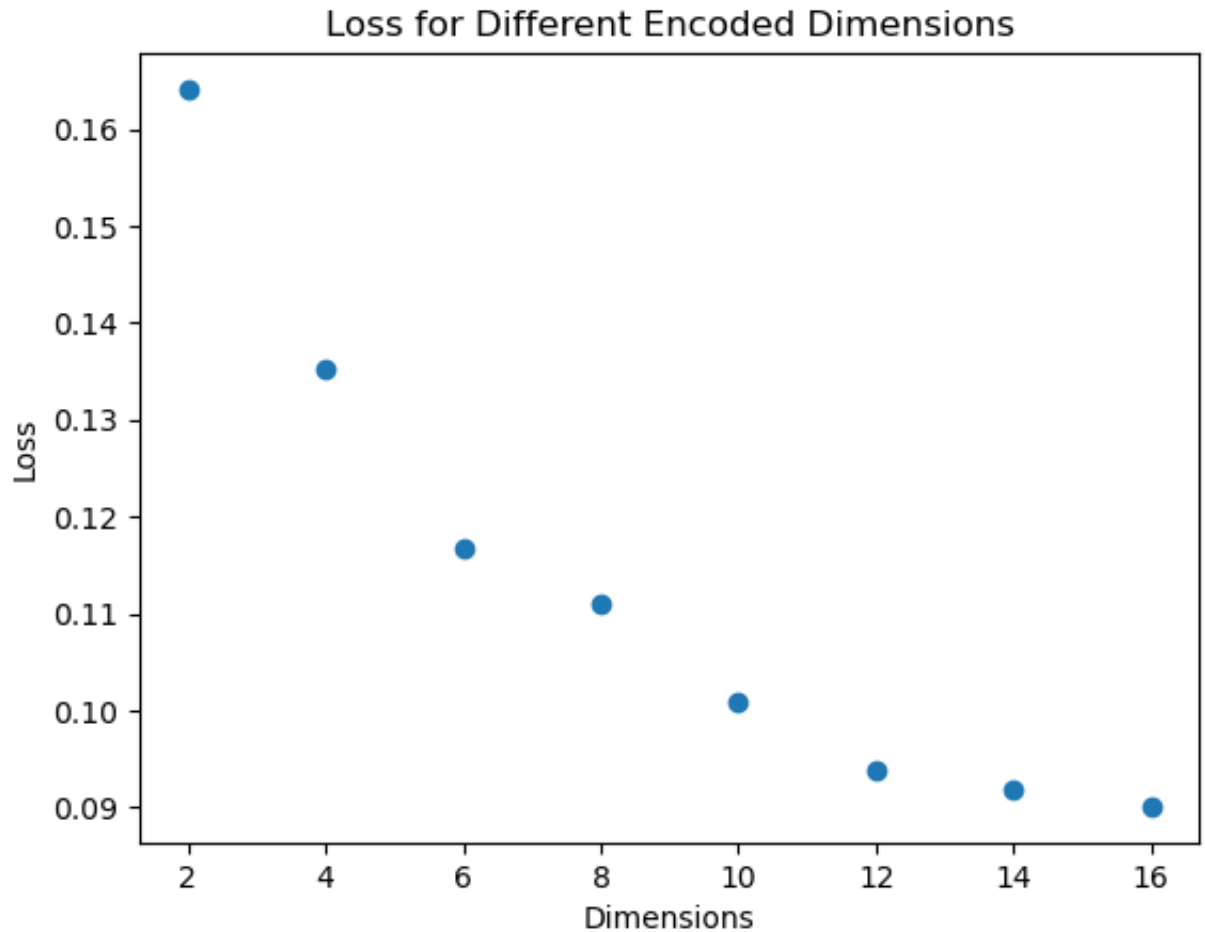
```
Dimensions, Loss, Accuracy: [2, 0.16411425173282623, 4, 0.13520190119
743347, 6, 0.11663084477186203, 8, 0.11094623804092407, 10, 0.1008804
589509964, 12, 0.09382828325033188, 14, 0.0918399766087532, 16, 0.090
```

08263051509857]

In [57]: *# Scatter Plot*

```
plt.scatter(dimensions, losses)
plt.xlabel('Dimensions')
plt.ylabel('Loss')
plt.title('Loss for Different Encoded Dimensions')
```

Out[57]: Text(0.5, 1.0, 'Loss for Different Encoded Dimensions')



2. **After** training an autoencoder with `encoding_dim=8`, apply noise (like the previous assignment) to *only* the input of the trained autoencoder (not the output). The output images should be without noise.

Print a few noisy images along with the output images to show they don't have noise.

In [42]: *# Building the model*

```
encoding_dim = 8

x = input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
# input_img_noise = input_img + np.random.normal(scale = 1.0, size = (
x = Dense(256, activation='relu')(x)
x = Dense(128, activation='relu')(x)
encoded = Dense(encoding_dim, activation='relu')(x)

# "decoded" is the lossy reconstruction of the input
x = Dense(128, activation='relu')(encoded)
x = Dense(256, activation='relu')(x)
decoded = Dense(784, activation='sigmoid')(x)

# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)

encoder = Model(input_img, encoded)

# create a placeholder for an encoded (32-dimensional) input
encoded_input = Input(shape=(encoding_dim,))
# retrieve the last layer of the autoencoder model
dcd1 = autoencoder.layers[-1]
dcd2 = autoencoder.layers[-2]
dcd3 = autoencoder.layers[-3]

# create the decoder model
decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))
```

In [43]: `np.shape(x_train), np.shape(x_test)`

Out[43]: ((60000, 784), (10000, 784))

In [50]: *# Noise and modeling*

```
x_train_noise = x_train + np.random.normal(scale = 0.5, size = (60000, 784))
x_test_noise = x_test + np.random.normal(scale = 0.5, size = (10000, 784))

autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

autoencoder.fit(x_train_noise, x_train,
                epochs=100,
                batch_size=256,
                shuffle=True,
                verbose=0,
                validation_data=(x_test_noise, x_test))
```

Out[50]: <keras.src.callbacks.History at 0x7f9267bcc460>

In [51]: *# Displaying images*

```
encoded_imgs = encoder.predict(x_test_noise)
decoded_imgs = decoder.predict(encoded_imgs)

n = 20 # how many digits we will display
plt.figure(figsize=(40, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test_noise[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```

313/313 [=====] - 0s 1ms/step

313/313 [=====] - 0s 1ms/step

